

DV0101EN-3-4-1-Waffle-Charts-Word-Clouds-and-Regression-Plots-py-v2.0

April 2, 2019

Waffle Charts, Word Clouds, and Regression Plots

0.1 Introduction

In this lab, we will learn how to create word clouds and waffle charts. Furthermore, we will start learning about additional visualization libraries that are based on Matplotlib, namely the library *seaborn*, and we will learn how to create regression plots using the *seaborn* library.

0.2 Table of Contents

1. Section ??
2. Section ??
3. Section ??
4. Section ??
5. Section ??
6. Section ??

1 Exploring Datasets with *pandas* and Matplotlib

Toolkits: The course heavily relies on *pandas* and **Numpy** for data wrangling, analysis, and visualization. The primary plotting library we will explore in the course is *Matplotlib*.

Dataset: Immigration to Canada from 1980 to 2013 - [International migration flows to and from selected countries - The 2015 revision](#) from United Nation's website

The dataset contains annual data on the flows of international migrants as recorded by the countries of destination. The data presents both inflows and outflows according to the place of birth, citizenship or place of previous / next residence both for foreigners and nationals. In this lab, we will focus on the Canadian Immigration data.

2 Downloading and Prepping Data

Import Primary Modules:

```
In [1]: import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library
from PIL import Image # converting images into arrays
```

Let's download and import our primary Canadian Immigration dataset using *pandas* `read_excel()` method. Normally, before we can do that, we would need to download a module which *pandas* requires to read in excel files. This module is **xlrd**. For your convenience, we have pre-installed this module, so you would not have to worry about that. Otherwise, you would need to run the following line of code to install the **xlrd** module:

```
!conda install -c anaconda xlrd --yes
```

Download the dataset and read it into a *pandas* dataframe:

```
In [2]: df_can = pd.read_excel('https://ibm.box.com/shared/static/lw190pt9zpy5bd1ptyg2aw15awomz9
                                sheet_name='Canada by Citizenship',
                                skiprows=range(20),
                                skipfooter=2)

                                print('Data downloaded and read into a dataframe!')
```

Data downloaded and read into a dataframe!

Let's take a look at the first five items in our dataset

```
In [3]: df_can.head()
```

```
Out[3]:
```

	Type	Coverage	OdName	AREA	AreaName	REG	\
0	Immigrants	Foreigners	Afghanistan	935	Asia	5501	
1	Immigrants	Foreigners	Albania	908	Europe	925	
2	Immigrants	Foreigners	Algeria	903	Africa	912	
3	Immigrants	Foreigners	American Samoa	909	Oceania	957	
4	Immigrants	Foreigners	Andorra	908	Europe	925	

	RegName	DEV	DevName	1980	...	2004	2005	2006	\
0	Southern Asia	902	Developing regions	16	...	2978	3436	3009	
1	Southern Europe	901	Developed regions	1	...	1450	1223	856	
2	Northern Africa	902	Developing regions	80	...	3616	3626	4807	
3	Polynesia	902	Developing regions	0	...	0	0	1	
4	Southern Europe	901	Developed regions	0	...	0	0	1	

	2007	2008	2009	2010	2011	2012	2013
0	2652	2111	1746	1758	2203	2635	2004
1	702	560	716	561	539	620	603
2	3623	4005	5393	4752	4325	3774	4331
3	0	0	0	0	0	0	0
4	1	0	0	0	0	1	1

[5 rows x 43 columns]

Let's find out how many entries there are in our dataset

```
In [4]: # print the dimensions of the dataframe
        print(df_can.shape)

(195, 43)
```

Clean up data. We will make some modifications to the original dataset to make it easier to create our visualizations. Refer to *Introduction to Matplotlib and Line Plots* and *Area Plots, Histograms, and Bar Plots* for a detailed description of this preprocessing.

```
In [5]: # clean up the dataset to remove unnecessary columns (eg. REG)
        df_can.drop(['AREA', 'REG', 'DEV', 'Type', 'Coverage'], axis = 1, inplace = True)

        # let's rename the columns so that they make sense
        df_can.rename (columns = {'OdName': 'Country', 'AreaName': 'Continent', 'RegName': 'Region'})

        # for sake of consistency, let's also make all column labels of type string
        df_can.columns = list(map(str, df_can.columns))

        # set the country name as index - useful for quickly looking up countries using .loc method
        df_can.set_index('Country', inplace = True)

        # add total column
        df_can['Total'] = df_can.sum (axis = 1)

        # years that we will be using in this lesson - useful for plotting later on
        years = list(map(str, range(1980, 2014)))
        print ('data dimensions:', df_can.shape)

data dimensions: (195, 38)
```

3 Visualizing Data using Matplotlib

Import matplotlib:

```
In [6]: %matplotlib inline

import matplotlib as mpl
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches # needed for waffle Charts

mpl.style.use('ggplot') # optional: for ggplot-like style

# check for latest version of Matplotlib
print ('Matplotlib version: ', mpl.__version__) # >= 2.0.0

Matplotlib version: 3.0.2
```

4 Waffle Charts

A waffle chart is an interesting visualization that is normally created to display progress toward goals. It is commonly an effective option when you are trying to add interesting visualization features to a visual that consists mainly of cells, such as an Excel dashboard.

Let's revisit the previous case study about Denmark, Norway, and Sweden.

```
In [7]: # let's create a new dataframe for these three countries
```

```
df_dsn = df_can.loc[['Denmark', 'Norway', 'Sweden'], :]
```

```
# let's take a look at our dataframe
```

```
df_dsn
```

```
Out[7]:
```

	Continent	Region	DevName	1980	1981	1982	1983	\
Country								
Denmark	Europe	Northern Europe	Developed regions	272	293	299	106	
Norway	Europe	Northern Europe	Developed regions	116	77	106	51	
Sweden	Europe	Northern Europe	Developed regions	281	308	222	176	

	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	\
Country				...								
Denmark	93	73	93	...	62	101	97	108	81	92	93	
Norway	31	54	56	...	57	53	73	66	75	46	49	
Sweden	128	158	187	...	205	139	193	165	167	159	134	

	2012	2013	Total
Country			
Denmark	94	81	3901
Norway	53	59	2327
Sweden	140	140	5866


```
[3 rows x 38 columns]
```

Unfortunately, unlike R, waffle charts are not built into any of the Python visualization libraries. Therefore, we will learn how to create them from scratch.

Step 1. The first step into creating a waffle chart is determining the proportion of each category with respect to the total.

```
In [8]: # compute the proportion of each category with respect to the total
```

```
total_values = sum(df_dsn['Total'])
```

```
category_proportions = [(float(value) / total_values) for value in df_dsn['Total']]
```

```
# print out proportions
```

```
for i, proportion in enumerate(category_proportions):
```

```
    print (df_dsn.index.values[i] + ': ' + str(proportion))
```

```
Denmark: 0.32255663965602777
```

```
Norway: 0.1924094592359848
```

```
Sweden: 0.48503390110798744
```

Step 2. The second step is defining the overall size of the waffle chart.

```
In [10]: width = 40 # width of chart
        height = 10 # height of chart

        total_num_tiles = width * height # total number of tiles

        print ('Total number of tiles is ', total_num_tiles)
```

Total number of tiles is 400

Step 3. The third step is using the proportion of each category to determine its respective number of tiles

```
In [11]: # compute the number of tiles for each category
        tiles_per_category = [round(proportion * total_num_tiles) for proportion in category_proportions]

        # print out number of tiles per category
        for i, tiles in enumerate(tiles_per_category):
            print (df_dsn.index.values[i] + ': ' + str(tiles))
```

Denmark: 129

Norway: 77

Sweden: 194

Based on the calculated proportions, Denmark will occupy 129 tiles of the waffle chart, Norway will occupy 77 tiles, and Sweden will occupy 194 tiles.

Step 4. The fourth step is creating a matrix that resembles the waffle chart and populating it.

```
In [12]: # initialize the waffle chart as an empty matrix
        waffle_chart = np.zeros((height, width))

        # define indices to loop through waffle chart
        category_index = 0
        tile_index = 0

        # populate the waffle chart
        for col in range(width):
            for row in range(height):
                tile_index += 1

                # if the number of tiles populated for the current category is equal to its count
                if tile_index > sum(tiles_per_category[0:category_index]):
                    # ...proceed to the next category
                    category_index += 1

                # set the class value to an integer, which increases with class
```

```
waffle_chart[row, col] = category_index

print ('Waffle chart populated!')

Waffle chart populated!
```

Let's take a peek at how the matrix looks like.

```
In [13]: waffle_chart
```

```
Out[13]: array([[1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.],
               [1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 2., 2., 2.,
                2., 2., 2., 2., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3., 3.,
                3., 3., 3., 3., 3., 3., 3., 3.]])
```

As expected, the matrix consists of three categories and the total number of each category's instances matches the total number of tiles allocated to each category.

Step 5. Map the waffle chart matrix into a visual.

```
In [14]: # instantiate a new figure object
fig = plt.figure()
```

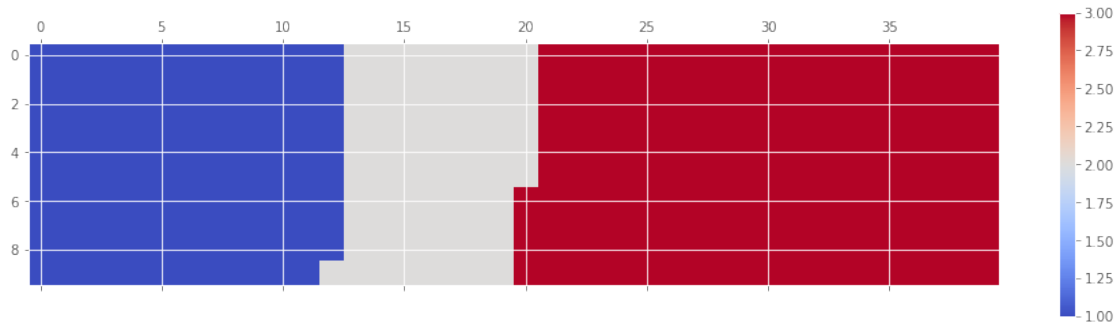
```

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

```

Out[14]: <matplotlib.colorbar.Colorbar at 0x7f65d0ad7e80>

<Figure size 432x288 with 0 Axes>



Step 6. Prettify the chart.

```

In [15]: # instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

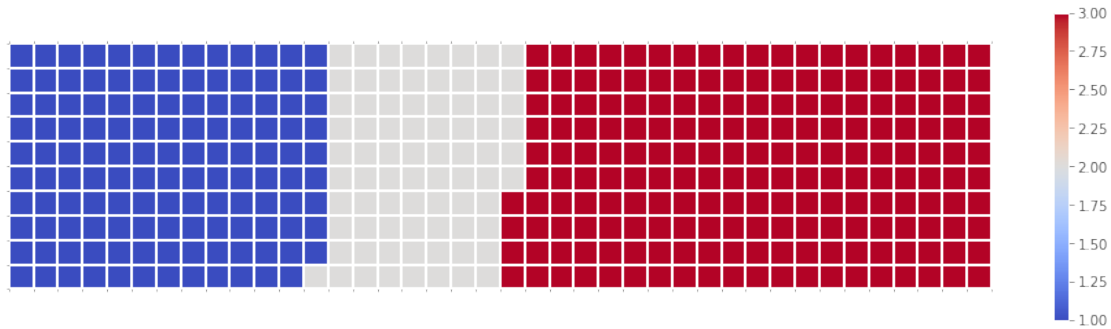
# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])

```

Out[15]: ([], <a list of 0 Text yticklabel objects>)

<Figure size 432x288 with 0 Axes>



Step 7. Create a legend and add it to chart.

```
In [16]: # instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])

# compute cumulative sum of individual categories to match color schemes between chart
values_cumsum = np.cumsum(df_dsn['Total'])
total_values = values_cumsum[len(values_cumsum) - 1]

# create legend
legend_handles = []
for i, category in enumerate(df_dsn.index.values):
    label_str = category + ' (' + str(df_dsn['Total'][i]) + ')'
    color_val = colormap(float(values_cumsum[i])/total_values)
    legend_handles.append(mpatches.Patch(color=color_val, label=label_str))

# add legend to chart
plt.legend(handles=legend_handles,
```




```

loc='lower center',
ncol=len(df_dsn.index.values),
bbox_to_anchor=(0., -0.2, 0.95, .1)
)

```

Out[16]: <matplotlib.legend.Legend at 0x7f65d1ddf358>

<Figure size 432x288 with 0 Axes>



And there you go! What a good looking *delicious* waffle chart, don't you think?

Now it would be very inefficient to repeat these seven steps every time we wish to create a waffle chart. So let's combine all seven steps into one function called *create_waffle_chart*. This function would take the following parameters as input:

1. **categories**: Unique categories or classes in dataframe.
2. **values**: Values corresponding to categories or classes.
3. **height**: Defined height of waffle chart.
4. **width**: Defined width of waffle chart.
5. **colormap**: Colormap class
6. **value_sign**: In order to make our function more generalizable, we will add this parameter to address signs that could be associated with a value such as %, \$, and so on. **value_sign** has a default value of empty string.

```

In [17]: def create_waffle_chart(categories, values, height, width, colormap, value_sign=''):

    # compute the proportion of each category with respect to the total
    total_values = sum(values)
    category_proportions = [(float(value) / total_values) for value in values]

    # compute the total number of tiles
    total_num_tiles = width * height # total number of tiles
    print ('Total number of tiles is', total_num_tiles)

    # compute the number of tiles for each category
    tiles_per_category = [round(proportion * total_num_tiles) for proportion in category

```

```

# print out number of tiles per category
for i, tiles in enumerate(tiles_per_category):
    print (df_dsn.index.values[i] + ': ' + str(tiles))

# initialize the waffle chart as an empty matrix
waffle_chart = np.zeros((height, width))

# define indices to loop through waffle chart
category_index = 0
tile_index = 0

# populate the waffle chart
for col in range(width):
    for row in range(height):
        tile_index += 1

        # if the number of tiles populated for the current category
        # is equal to its corresponding allocated tiles...
        if tile_index > sum(tiles_per_category[0:category_index]):
            # ...proceed to the next category
            category_index += 1

        # set the class value to an integer, which increases with class
        waffle_chart[row, col] = category_index

# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='-', linewidth=2)

plt.xticks([])
plt.yticks([])

# compute cumulative sum of individual categories to match color schemes between ch

```

```

values_cumsum = np.cumsum(values)
total_values = values_cumsum[len(values_cumsum) - 1]

# create legend
legend_handles = []
for i, category in enumerate(categories):
    if value_sign == '%':
        label_str = category + ' (' + str(values[i]) + value_sign + ')'
    else:
        label_str = category + ' (' + value_sign + str(values[i]) + ')'

    color_val = colormap(float(values_cumsum[i])/total_values)
    legend_handles.append(mpatches.Patch(color=color_val, label=label_str))

# add legend to chart
plt.legend(
    handles=legend_handles,
    loc='lower center',
    ncol=len(categories),
    bbox_to_anchor=(0., -0.2, 0.95, .1)
)

```

Now to create a waffle chart, all we have to do is call the function `create_waffle_chart`. Let's define the input parameters:

```

In [18]: width = 40 # width of chart
        height = 10 # height of chart

        categories = df_dsn.index.values # categories
        values = df_dsn['Total'] # corresponding values of categories

        colormap = plt.cm.coolwarm # color map class

```

And now let's call our function to create a waffle chart.

```

In [19]: create_waffle_chart(categories, values, height, width, colormap)

```

```

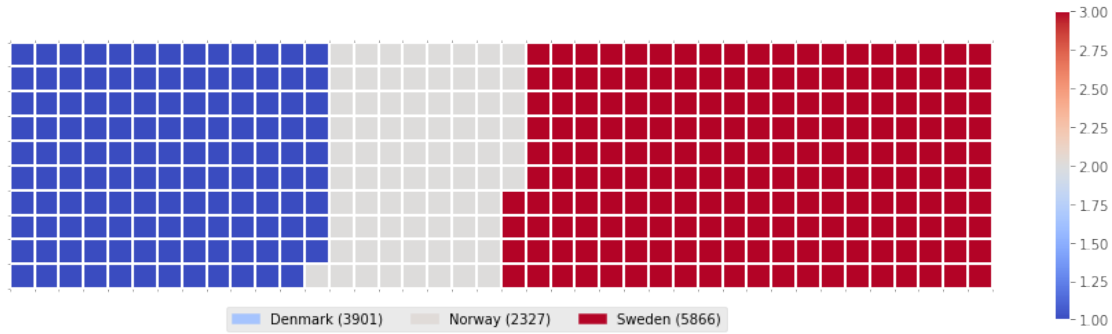
Total number of tiles is 400
Denmark: 129
Norway: 77
Sweden: 194

```

```

<Figure size 432x288 with 0 Axes>

```



There seems to be a new Python package for generating waffle charts called [PyWaffle](#), but it looks like the repository is still being built. But feel free to check it out and play with it.

5 Word Clouds

Word clouds (also known as text clouds or tag clouds) work in a simple way: the more a specific word appears in a source of textual data (such as a speech, blog post, or database), the bigger and bolder it appears in the word cloud.

Luckily, a Python package already exists in Python for generating word clouds. The package, called `wordcloud` was developed by **Andreas Mueller**. You can learn more about the package by following this [link](#).

Let's use this package to learn how to generate a word cloud for a given text document. First, let's install the package.

```
In [20]: # install wordcloud
         !conda install -c conda-forge wordcloud==1.4.1 --yes

         # import package and its set of stopwords
         from wordcloud import WordCloud, STOPWORDS

         print ('Wordcloud is installed and imported!')
```

Collecting package metadata: done

Solving environment: /

The environment is inconsistent, please check the package plan carefully

The following packages are causing the inconsistency:

- defaults/linux-64::anaconda==5.3.1=py37_0
- defaults/linux-64::astropy==3.0.4=py37h14c3975_0
- defaults/linux-64::bkcharts==0.2=py37_0
- defaults/linux-64::blaze==0.11.3=py37_0
- defaults/linux-64::bokeh==0.13.0=py37_0
- defaults/linux-64::bottleneck==1.2.1=py37h035aef0_1
- defaults/linux-64::dask==0.19.1=py37_0
- defaults/linux-64::datashape==0.5.4=py37_1

```

- defaults/linux-64::mkl-service==1.1.2=py37h90e4bf4_5
- defaults/linux-64::numba==0.39.0=py37h04863e7_0
- defaults/linux-64::numexpr==2.6.8=py37hd89afb7_0
- defaults/linux-64::odo==0.5.1=py37_0
- defaults/linux-64::pytables==3.4.4=py37ha205bf6_0
- defaults/linux-64::pytest-arraydiff==0.2=py37h39e3cac_0
- defaults/linux-64::pytest-astropy==0.4.0=py37_0
- defaults/linux-64::pytest-doctestplus==0.1.3=py37_0
- defaults/linux-64::pywavelets==1.0.0=py37hdd07704_0
- defaults/linux-64::scikit-image==0.14.0=py37hf484d3e_1
done

```

Package Plan

environment location: /home/jupyterlab/conda

added / updated specs:

```
- wordcloud==1.4.1
```

The following packages will be downloaded:

package	build		
wordcloud-1.4.1	py36_0	324 KB	conda-forge
Total:		324 KB	

The following NEW packages will be INSTALLED:

```
wordcloud          conda-forge/linux-64::wordcloud-1.4.1-py36_0
```

Downloading and Extracting Packages

```
wordcloud-1.4.1      | 324 KB      | ##### | 100%
```

Preparing transaction: done

Verifying transaction: done

Executing transaction: done

Wordcloud is installed and imported!

Word clouds are commonly used to perform high-level analysis and visualization of text data. Accordingly, let's digress from the immigration dataset and work with an example that involves analyzing text data. Let's try to analyze a short novel written by **Lewis Carroll** titled *Alice's Adventures in Wonderland*. Let's go ahead and download a *.txt* file of the novel.

```
In [25]: # download file and save as alice_novel.txt
!wget --quiet https://ibm.box.com/shared/static/m54sjtrshpt5su20dzesl5en9xa5vfz1.txt -C
```

File downloaded and saved!

```
In [26]: stopwords = set(STOPWORDS)
```

```
# generate the word cloud
alice_wc.generate(alice_novel)
```

```
In [28]: # display the word cloud
plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```




```
plt.axis('off')
plt.show()
```



Excellent! This looks really interesting! Another cool thing you can implement with the `word_cloud` package is superimposing the words onto a mask of any shape. Let's use a mask of Alice and her rabbit. We already created the mask for you, so let's go ahead and download it and call it *alice_mask.png*.

```
In [31]: # download image
!wget --quiet https://ibm.box.com/shared/static/3mpxgaf6muer6af7t1nvqkw9cqj85ibm.png -C

# save mask to alice_mask
alice_mask = np.array(Image.open('alice_mask.png'))

print('Image downloaded and saved!')
```

Image downloaded and saved!

Let's take a look at how the mask looks like.

```
In [32]: fig = plt.figure()
fig.set_figwidth(14) # set width
fig.set_figheight(18) # set height

plt.imshow(alice_mask, cmap=plt.cm.gray, interpolation='bilinear')
plt.axis('off')
plt.show()
```




Shaping the word cloud according to the mask is straightforward using word_cloud package. For simplicity, we will continue using the first 2000 words in the novel.

```
In [33]: # instantiate a word cloud object
         alice_wc = WordCloud(background_color='white', max_words=2000, mask=alice_mask, stopwords=stopwords)

         # generate the word cloud
         alice_wc.generate(alice_novel)

         # display the word cloud
         fig = plt.figure()
         fig.set_figwidth(14) # set width
         fig.set_figheight(18) # set height

         plt.imshow(alice_wc, interpolation='bilinear')
```

```
plt.axis('off')
plt.show()
```



Really impressive!

Unfortunately, our immigration data does not have any text data, but where there is a will there is a way. Let's generate sample text data from our immigration dataset, say text data of 90 words.

Let's recall how our data looks like.

```
In [34]: df_can.head()
```

Country	Continent	Region	DevName	1980	1981	\
Afghanistan	Asia	Southern Asia	Developing regions	16	39	
Albania	Europe	Southern Europe	Developed regions	1	0	

Algeria	Africa	Northern Africa	Developing regions	80	67
American Samoa	Oceania	Polynesia	Developing regions	0	1
Andorra	Europe	Southern Europe	Developed regions	0	0

	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	\
Country						...					
Afghanistan	39	47	71	340	496	...	3436	3009	2652	2111	
Albania	0	0	0	0	1	...	1223	856	702	560	
Algeria	71	69	63	44	69	...	3626	4807	3623	4005	
American Samoa	0	0	0	0	0	...	0	1	0	0	
Andorra	0	0	0	0	2	...	0	1	1	0	

	2009	2010	2011	2012	2013	Total
Country						
Afghanistan	1746	1758	2203	2635	2004	58639
Albania	716	561	539	620	603	15699
Algeria	5393	4752	4325	3774	4331	69439
American Samoa	0	0	0	0	0	6
Andorra	0	0	0	1	1	15

[5 rows x 38 columns]

And what was the total immigration from 1980 to 2013?

```
In [35]: total_immigration = df_can['Total'].sum()
total_immigration
```

```
Out[35]: 6409153
```

Using countries with single-word names, let's duplicate each country's name based on how much they contribute to the total immigration.

```
In [36]: max_words = 90
word_string = ''
for country in df_can.index.values:
    # check if country's name is a single-word name
    if len(country.split(' ')) == 1:
        repeat_num_times = int(df_can.loc[country, 'Total']/float(total_immigration)*max_words)
        word_string = word_string + ((country + ' ') * repeat_num_times)

# display the generated text
word_string
```

```
Out[36]: 'China China China China China China China China China Colombia Egypt France Guyana Hai
```

We are not dealing with any stopwords here, so there is no need to pass them when creating the word cloud.

```
In [37]: # create the word cloud
wordcloud = WordCloud(background_color='white').generate(word_string)

print('Word cloud created!')
```

Word cloud created!

```
In [38]: # display the cloud
fig = plt.figure()
fig.set_figwidth(14)
fig.set_figheight(18)

plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()
```



According to the above word cloud, it looks like the majority of the people who immigrated came from one of 15 countries that are displayed by the word cloud. One cool visual that you could build, is perhaps using the map of Canada and a mask and superimposing the word cloud on top of the map of Canada. That would be an interesting visual to build!

6 Regression Plots

Seaborn is a Python visualization library based on matplotlib. It provides a high-level interface for drawing attractive statistical graphics. You can learn more about *seaborn* by following this [link](#) and more about *seaborn* regression plots by following this [link](#).

In lab *Pie Charts, Box Plots, Scatter Plots, and Bubble Plots*, we learned how to create a scatter plot and then fit a regression line. It took ~20 lines of code to create the scatter plot along with the regression fit. In this final section, we will explore *seaborn* and see how efficient it is to create regression lines and fits using this library!

Let's first install *seaborn*

```
In [39]: # install seaborn
!pip install seaborn

# import library
import seaborn as sns

print('Seaborn installed and imported!')
```

```
Requirement already satisfied: seaborn in /home/jupyterlab/conda/lib/python3.6/site-packages (0.
Requirement already satisfied: scipy>=0.14.0 in /home/jupyterlab/conda/lib/python3.6/site-package
Requirement already satisfied: numpy>=1.9.3 in /home/jupyterlab/conda/lib/python3.6/site-package
Requirement already satisfied: matplotlib>=1.4.3 in /home/jupyterlab/conda/lib/python3.6/site-pa
Requirement already satisfied: pandas>=0.15.2 in /home/jupyterlab/conda/lib/python3.6/site-packa
Requirement already satisfied: cyclor>=0.10 in /home/jupyterlab/conda/lib/python3.6/site-package
Requirement already satisfied: kiwisolver>=1.0.1 in /home/jupyterlab/conda/lib/python3.6/site-pa
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /home/jupyterlab/cond
Requirement already satisfied: python-dateutil>=2.1 in /home/jupyterlab/conda/lib/python3.6/site
Requirement already satisfied: pytz>=2011k in /home/jupyterlab/conda/lib/python3.6/site-packages
Requirement already satisfied: six in /home/jupyterlab/conda/lib/python3.6/site-packages (from c
Requirement already satisfied: setuptools in /home/jupyterlab/conda/lib/python3.6/site-packages
Seaborn installed and imported!
```

Create a new dataframe that stores that total number of landed immigrants to Canada per year from 1980 to 2013.

```
In [40]: # we can use the sum() method to get the total population per year
df_tot = pd.DataFrame(df_can[years].sum(axis=0))

# change the years to type float (useful for regression later on)
df_tot.index = map(float,df_tot.index)

# reset the index to put in back in as a column in the df_tot dataframe
df_tot.reset_index(inplace = True)

# rename columns
df_tot.columns = ['year', 'total']

# view the final dataframe
df_tot.head()
```

```
Out[40]:
```


	year	total
0	1980.0	99137

```
1 1981.0 110563
2 1982.0 104271
3 1983.0 75550
4 1984.0 73417
```

With *seaborn*, generating a regression plot is as simple as calling the **regplot** function.


```
In [41]: import seaborn as sns
         ax = sns.regplot(x='year', y='total', data=df_tot)
```

```
/home/jupyterlab/conda/lib/python3.6/site-packages/scipy/stats/stats.py:1713: FutureWarning: Use
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```




This is not magic; it is *seaborn*! You can also customize the color of the scatter plot and regression line. Let's change the color to green.

```
In [42]: import seaborn as sns
         ax = sns.regplot(x='year', y='total', data=df_tot, color='green')
```



You can always customize the marker shape, so instead of circular markers, let's use '+'.

```
In [43]: import seaborn as sns
ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+')
```



Let's blow up the plot a little bit so that it is more appealing to the sight.

```
In [44]: plt.figure(figsize=(15, 10))
         ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+')
```



And let's increase the size of markers so they match the new size of the figure, and add a title and x- and y-labels.

```
In [45]: plt.figure(figsize=(15, 10))
         ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+', scatter_kws={
            's': 100})

         ax.set(xlabel='Year', ylabel='Total Immigration') # add x- and y-labels
         ax.set_title('Total Immigration to Canada from 1980 - 2013') # add title
```

```
Out[45]: Text(0.5, 1.0, 'Total Immigration to Canada from 1980 - 2013')
```





And finally increase the font size of the tickmark labels, the title, and the x- and y-labels so they don't feel left out!

```
In [46]: plt.figure(figsize=(15, 10))
```

```
sns.set(font_scale=1.5)
```

```
ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+', scatter_k  
ax.set(xlabel='Year', ylabel='Total Immigration')  
ax.set_title('Total Immigration to Canada from 1980 - 2013')
```

```
Out[46]: Text(0.5, 1.0, 'Total Immigration to Canada from 1980 - 2013')
```



Amazing! A complete scatter plot with a regression fit with 5 lines of code only. Isn't this really amazing?


If you are not a big fan of the purple background, you can easily change the style to a white plain background.

```
In [47]: plt.figure(figsize=(15, 10))
```

```
sns.set(font_scale=1.5)
sns.set_style('ticks') # change background to white background
```

```
ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+', scatter_k
ax.set(xlabel='Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
```

```
Out[47]: Text(0.5, 1.0, 'Total Immigration to Canada from 1980 - 2013')
```




Or to a white background with gridlines.

```
In [48]: plt.figure(figsize=(15, 10))
```

```
sns.set(font_scale=1.5)
sns.set_style('whitegrid')
```

```
ax = sns.regplot(x='year', y='total', data=df_tot, color='green', marker='+', scatter_k
ax.set(xlabel='Year', ylabel='Total Immigration')
ax.set_title('Total Immigration to Canada from 1980 - 2013')
```

```
Out[48]: Text(0.5, 1.0, 'Total Immigration to Canada from 1980 - 2013')
```



Question: Use seaborn to create a scatter plot with a regression line to visualize the total immigration from Denmark, Sweden, and Norway to Canada from 1980 to 2013.

In []: *### type your answer here*

In []:

Double-click [here](#) for the solution.
 Double-click [here](#) for the solution.
 Double-click [here](#) for the solution.

6.0.1 Thank you for completing this lab!

This notebook was created by [Alex Aklson](#). I hope you found this lab interesting and educational. Feel free to contact me if you have any questions!

This notebook is part of a course on **Coursera** called *Data Visualization with Python*. If you accessed this notebook outside the course, you can take this course online by clicking [here](#).

Copyright © 2018 [Cognitive Class](#). This notebook and its source code are released under the terms of the [MIT License](#).